

Language Modeling and Smoothing

Deadline: 17th September

Submission format: Upload in your previous github repo in a separate directory.

1. Download any of these text books from Project Gutenberg
 - a. Alice in Wonderland: [Alice's Adventures in Wonderland](#)
 - b. Sherlock Holmes: [The Adventures of Sherlock Holmes](#)
 2. Parse the dataset into sentences using sentence [tokenizer](#) and divide it into 80/20 ratio. Keep 80% dataset for training N-grams and keep 20% for test. You can filter out unnecessary symbols, newlines, etc. You can add symbols <s> and </s> to mark sentence start and end.
 3. Compute MLE for unigram, bigram, trigrams and quadgrams. How many n-grams are possible and how many actually exists?
 4. Develop a system that has two functions:
 - a. Generator(model_name): generates sentences by utilizing MLEs from specified n-gram model. Sampling from multinomial distribution can be done using a predefined [function](#). Note, 5-10 sentences would suffice for this task.
 - b. Probability(sentence,model_name): Compute the probability of a given sentence in log-space. Note you can provide any sentence, however, a random sentence will mostly lead to zero probability. The better idea is to take sentences from the corpus itself.
 5. Implement add-1 smoothing for bigram model and give 2-3 examples where drastic change in the count occurs post-smoothing. Can you explain this drastic change in a sentence?
 6. Do you observe the constant discounting value 'd' by implementing Good-turing smoothing technique? If yes, what is the value of 'd'?
- Hint:** You can check for n-grams having original counts between 1-10.

7. Compute the perplexity value for the test dataset for the bigram model using add-1 and Good-turing. Which performs better?